

Masterarbeit im Fach Informatik  
RHEINISCH-WESTFÄLISCHE TECHNISCHE HOCHSCHULE AACHEN  
Lehrstuhl für Informatik 6  
Prof. Dr.-Ing. H. Ney

---

# Unsupervised Learning of Neural Network Lexicon and Cross-lingual Word Embedding

---

20. September 2018

vorgelegt von:  
Autor Jiahui Geng  
Matrikelnummer 365655

Gutachter:  
Prof. Dr.-Ing. H. Ney  
Prof. B. Leibe, Ph. D.

Betreuer:  
M.Sc. Yunsu Kim



# Erklärung

Geng, Jiahui

Name, Vorname

365655

Matrikelnummer (freiwillige Angabe)

Ich versichere hiermit an Eides Statt, dass ich die vorliegende ~~Arbeit~~/Bachelorarbeit/  
~~Masterarbeit~~\* mit dem Titel

Unsupervised Learning of Neural Network Lexicon and Cross-lingual Word Embedding

selbständig und ohne unzulässige fremde Hilfe erbracht habe. Ich habe keine anderen als die angegebenen Quellen und Hilfsmittel benutzt. Für den Fall, dass die Arbeit zusätzlich auf einem Datenträger eingereicht wird, erkläre ich, dass die schriftliche und die elektronische Form vollständig übereinstimmen. Die Arbeit hat in gleicher oder ähnlicher Form noch keiner Prüfungsbehörde vorgelegen.

Aachen, 20. September 2018

Ort, Datum

Unterschrift

\*Nichtzutreffendes bitte streichen

## Belehrung:

### § 156 StGB: Falsche Versicherung an Eides Statt

Wer vor einer zur Abnahme einer Versicherung an Eides Statt zuständigen Behörde eine solche Versicherung falsch abgibt oder unter Berufung auf eine solche Versicherung falsch aussagt, wird mit Freiheitsstrafe bis zu drei Jahren oder mit Geldstrafe bestraft.

### § 161 StGB: Fahrlässiger Falscheid; fahrlässige falsche Versicherung an Eides Statt

(1) Wenn eine der in den §§ 154 bis 156 bezeichneten Handlungen aus Fahrlässigkeit begangen worden ist, so tritt Freiheitsstrafe bis zu einem Jahr oder Geldstrafe ein.

(2) Straflosigkeit tritt ein, wenn der Täter die falsche Angabe rechtzeitig berichtigt. Die Vorschriften des § 158 Abs. 2 und 3 gelten entsprechend.

Die vorstehende Belehrung habe ich zur Kenntnis genommen:

Aachen, 20. September 2018

Ort, Datum

Unterschrift



# Abstract

The unsupervised learning of cross-lingual word embedding offers elegant word matches between languages, but there are fundamental limitations in sentence translation. In this work, we propose a simple yet effective method to improve the word-by-word sentence translation based on cross-lingual word embedding, using only monolingual corpora but without any back-translation. We integrate a language model for context-aware beam search, we also combine with a novel denoising autoencoder to handle reordering. Our system surpasses state-of-the-art unsupervised machine translation system without costly iterative training. We also analyze the effects of parameters and artificial designed features in our model. We also propose a novel data-driven unsupervised learning of cross-lingual word embedding, which can be actually describe as an iterative process combing the translation process and SGD training of word embedding. Since the quality of cross-lingual word embedding effects the performance of our model. Our contributions to embedding learning and translation system can provide better understanding of learning the cross-lingual word embedding and its usage in translation.



# Contents

<b>Abstract</b>	<b>v</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Related Work . . . . .	2
1.1.1 Word Embedding . . . . .	2
1.1.2 Unsupervised Machine Translation . . . . .	3
1.2 Outline . . . . .	4
1.3 Notation . . . . .	4
<b>2 Machine Translation</b>	<b>7</b>
2.1 Statistical Machine Translation . . . . .	7
2.1.1 Word-Based Model . . . . .	7
2.1.2 Phrase-Based Model . . . . .	10
2.2 Neural Machine Translation . . . . .	10
2.2.1 Seq2seq Model . . . . .	11
2.2.2 Attention Mechanism . . . . .	12
2.2.3 Transformer . . . . .	15
2.3 Unsupervised Machine Translation . . . . .	18
2.3.1 Decipherment . . . . .	18
2.3.2 Neural Unsupervised Machine Translation . . . . .	19
<b>3 Word Embedding</b>	<b>23</b>
3.1 Monolingual Embedding . . . . .	23
3.1.1 CBOW and Skip-gram Model . . . . .	23
3.1.2 FastText . . . . .	25
3.2 Cross-lingual Word Embedding . . . . .	25
3.2.1 *** . . . . .	26
3.2.2 Orthogonal Constraints . . . . .	27
3.2.3 CSLS Loss . . . . .	28
<b>4 Cross-lingual Word Embedding without Parallel Data</b>	<b>29</b>
4.1 Initialization . . . . .	29
4.1.1 Iterative Closest Point Method . . . . .	29
4.1.2 Adversarial training . . . . .	30
4.2 Nearest Neighbor Search . . . . .	31

4.3	Iterative Training . . . . .	32
4.3.1	Self-learning with Procrustes . . . . .	33
4.3.2	Data-driven Method . . . . .	33
<b>5</b>	<b>Sentence Translation</b>	<b>35</b>
5.1	Context-aware Beam Search . . . . .	35
5.1.1	Language Model . . . . .	35
5.1.2	Beam Search . . . . .	36
5.2	Denoising Neural Network . . . . .	36
5.2.1	Denoising Auto-encoder . . . . .	36
5.2.2	Noise Model . . . . .	37
<b>6</b>	<b>Experiments</b>	<b>41</b>
6.1	Corpus Statistics . . . . .	41
6.2	title . . . . .	41
6.3	title . . . . .	41
<b>7</b>	<b>Conclusion</b>	<b>45</b>
<b>A</b>	<b>Appendix</b>	<b>47</b>
	<b>List of Figures</b>	<b>49</b>
	<b>List of Tables</b>	<b>51</b>
	<b>Bibliography</b>	<b>53</b>



# Chapter 1

## Introduction

Neural machine translation (NMT) has recently become the dominant method to machine translation task. In comparison to the traditional statistical machine translation (SMT), NMT systems are trained end-to-end, taking advantages of the continuous representation of the hidden states that greatly alleviate the sparsity problem and make better use of more contextual information. Building a good machine translation system requires a huge collections of parallel data. NMT systems often fail when the training data is not enough. The lack of parallel corpora is actually a problem. Parallel corpora are costly to build because they requires huge human labor and time and sometime expertise of corresponding languages. Conversely, the monolingual corpora is readily available.

Several approaches has been proposed to alleviate this issue, for instance, triangulation and semi-supervised learning techniques, how these systems still need a strong cross-lingual signal. Unsupervised machine translation uses only monolingual data of the source and target language to train the model.

Recently the work on cross-lingual word embedding can help to eliminate the need of cross-lingual information. Cross-lingual word embedding can be defined as the word embedding of multiple languages in a joint embedding space. In our work, we consider the learning of cross-lingual word embedding space as the task of learning the mapping from source embedding space to target embedding space.

In this work, we propose a simple yet efficient unsupervised machine translation system based on cross-lingual word embedding, we build the system by combining the context-aware beam search and using denoising autoencoder to handle reordering. Our system surpasses the state-of-the-art performance of unsupervised machine translation systems without iterative training. We also analyze the effect of detailed model parameters. We analyze the factors for cross-lingual word embedding training, propose a novel unsupervised training method using stochastic gradient descent instead of popular Procrustes training. We demonstrate that our algorithm can achieve the same performance with the state-of-the-art unsupervised methods.

For future work we hope to implement an iterative algorithm to polish our unsupervised machine translation model

## 1.1 Related Work

### 1.1.1 Word Embedding

Traditional language processing systems treat words as discrete atomic symbols, they assign for each word a specific id number. Such encodings are arbitrary, and it does not provide any information about the relations that may exist between individual symbols. In comparison, with learning algorithm each word can be represented in high dimensional continuous space. According to the distributional hypothesis, similar words tends to occurs with similar neighbors, so similar words have similar word representations. Mikolov et al. [2013b] first noticed that continuous embedding spaces exhibit similar structures across languages, even when for distant language pairs for example English-Vietnamese.

The series of researches about word embedding started from Mikolov et al. [2013a]. In this work, he just proposed two famous structures: continuous bag-of-words model (CBOW) and the skip-gram model to learn the word embedding and such embeddings have good properties for measure the syntactic and semantic word similarities and in his further work: Mikolov et al. [2013c], he discussed about learning the embedding for phrases. Pennington et al. [2014] proposed a regression model, which combines the global matrix factorization and local context window methods, make the training process more interpretable. Bojanowski et al. [2016] tried to represent each word as a bag of character  $n$ -grams, he assigned for each character  $n$ -gram a distinct embedding, in this way, we can exploit the subword information. Mikolov et al. [2013b] first noticed that continuous word embedding space exhibits similar structures according languages, even for distant language pairs, which is also the start work for cross-lingual word embedding, and this works also inspired the work for unsupervised machine translation, since they can provide heuristic information. In detail they tried to learn similarity by learning a linear mapping from a source to a target embedding space. They employed a parallel vocabulary of five thousand words as anchor points to learn this mapping and evaluated their approach on a word translation task. Xing et al. [2015] showed that results can be improved by enforce an orthogonal constraint on the linear mapping.

Recently, Artetxe et al. [2017a] proposed an iterative method that align the word embedding spaces gradually. Thanks to his contribution, the vocabulary size for learning embeddings can be reduced to 25 word pairs and even with Arabic numerals. The performance is comparable to the learning with large vocabulary dictionary Cao et al. [2016] proposed a distribution-based model, which is a modified CBOW model which tries to minimization between the distribution dissimilarity between source and target embeddings. Here the distribution refers the mean and variance. Zhang et al. [2017] proposed a method using adversarial training without any parallel data. The discriminator tried to discriminate if the embedding from the source

side and the generator aimed to learning the mapping from source embedding space to target one. These methods are totally unsupervised but the performances are far worse than the supervised training. Conneau et al. [2017] simplifies the adversarial training structure with different loss function for generator and discriminator. The performance got much improved, in addition, he viewed the word translation task as a word retrieval task and find the nearest neighbor searching suffers from the hubness problem. They further proposed to use cross-domain similarity local scaling (CSLS) penalize the hubs and a validation criterion for unsupervised model selection.

### 1.1.2 Unsupervised Machine Translation

As mentioned previously, the lack of parallel corpora motivate people to use monolingual data to improve the machine translation system. Some researchers tried to use triangulation techniques Cohn and Lapata [2007] and semi-supervised approaches Cheng et al. [2016] to address. But these methods still require parallel corpora.

As to total unsupervised machine translation, which only takes advantages of monolingual data, Ravi and Knight [2011] first proposed some ideas about the unsupervised machine model, he considered the unsupervised translation as a deciphering problem, where the source language is considered as ciphertext. They also proposed iterative EM method and Bayesian decipherment to train this unsupervised model. To solve the bottleneck of such model, mainly the huge memory is required to store the candidates search space, Based on deciphering model, Nuhn et al. [2012] tried to limit search candidates according to the word similarity. Nuhn and Ney [2014] limited the search space by using beam search and preselection search. Kim et al. [2017] enforced the sparsity with a simple threshold for the lexicon. He also tried to initialize the lexicon training with word classes, which efficiently boosts the performance. Although initially not based on distributional semantics, recent studies show that the use of word embeddings can bring significant improvement in statistical decipherment Duong et al. [2016].

He et al. [2016] made an important contribution to train neural model for unsupervised machine translation. More concretely, their method trains two agents to translate in opposite directions (e.g. French  $\rightarrow$  English and English  $\rightarrow$  French) and make them teach each other through a reinforcement learning process. While promising, this approach still requires a small parallel corpora for a warm start. Artetxe et al. [2017b] Lample et al. [2017] proposed two bi-directional unsupervised machine translation model which totally rely only on monolingual corpora in each language. These two models both need to use cross-lingual word embedding to initialize the MT system and, train the sequence-to-sequence system as denoising autoencoder and turn the unsupervised training into a supervised one by introduc-

ing back-translation techniques. The most important principle for these two work are the sharing embedding space. Sentences from different languages are encoded into a shared embedding spaces and then translated into specific language with different decoder. Lample et al. [2018] proposed two model variants, a neural and a phrase-based model. These models have fewer hyper-parameters. And these models performs well also on some distant and low-resource languages.

## 1.2 Outline

In summary, this work makes the following main contributions:

The remainder of this thesis is structured as follows. In Section 1.3 we introduce Chapter 2 introduces the development of machine translation systems from statistical models to to neural models. Some basic techniques and principles for machine translation are also included. Chapter 3 gives more information, I do a survey of training and applications details of cross-lingual word embedding. We discuss the unsupervised learning of cross-lingual word embedding and our efforts on data-drive training. We describe our unsupervised machine translation model in chapter 5, which contains mainly context-aware part: with the help of language model and denoising autoencoder aimed at reordering. We demonstrate the results of cross-lingual word embedding model and unsupervised machine model. We will compare the performance with the state-of-the-art model. In Chapter 7 we summarize our work which helps better understanding of learning the corss-lingual word embedding and its usage in translation domain.

## 1.3 Notation

Normally we denote

- source sentence  $f_1^J := f_1 \cdots f_j \cdots f_J$
- target sentence  $e_1^I := e_1 \cdots e_i \cdots e_I$
- single character  $e, f$  separately
- sentences  $e_1^N, f_1^N$  when word-by-word translating
- $e$  and  $f$  the source and target word embedding
- $E, F$  are the corresponding embedding matrices
- source and target Corpora  $\mathcal{S}, \mathcal{T}$

- model probability with subscripts  $s, t$  such as:  
 $P_{s \rightarrow s}(\cdots), P_{t \rightarrow t}(\cdots)$  for denoising autoencoder,  $P_{s \rightarrow t}(\cdots), P_{t \rightarrow s}(\cdots)$  for translation model



## Chapter 2

# Machine Translation

This chapter describes the classical or start-of-the-art machine translation models from statistical machine translation model to neural machine model. The relative research works on unsupervised machine translation will also be included.

### 2.1 Statistical Machine Translation

Statistical machine translation has achieved success until the beginning of this century. The initial statistical models for machine translation are based on words as atomic units that may be translated, inserted, dropped and reordered. In statistical machine translation, we use both a translation model and a language model which ensures fluent output. Later the statistical machine translation prefers to use translation of phrases as atomic units. These phrases are any contiguous sequences of words, not necessarily linguistic entities. In this approach, the input sentence is broken up into a sequence of phrases, these phrases are mapped one-to-one to output phrases, which may be reordered.

#### 2.1.1 Word-Based Model

**Noisy-channel model:** Noisy channel model based on the notion of a noisy channel from Shannon's information theory has been applied to many language processing problems. Assume the source sentence is a distorted message omitted from the target sentence, we have a model on how the message is distorted (translation model  $Pr(f_1^J|e_1^I)$ ) and also a model on which original messages are probable (language model  $Pr(e_1^I)$ ), our task is to find the best translation  $e_1^I$  for an input foreign sentence.

$$\begin{aligned}\operatorname{argmax}_{e_1^I}\{Pr(e_1^I|f_1^J)\} &= \operatorname{argmax}_{e_1^I}\frac{Pr(f_1^J|e_1^I)Pr(e_1^I)}{Pr(f_1^J)} \\ &= \operatorname{argmax}_{e_1^I}\{Pr(f_1^J|e_1^I) \cdot Pr(e_1^I)\}\end{aligned}$$

The basic model for word-based model is noisy-channel model. The task of word alignment is an artifact of word-based machine translation. Alignment models target at the reordering problem for word-based translation.

**Alignment model:** The position  $j$  in the source sentence is aligned with the position  $i$  in the target sentence when translating, denoted as  $i = a_j$ . Alignment model is a global reordering model. For whole sentence, we denote the alignment as

$$a_1^J := a_1 \cdots a_J$$

### Zero-Order Models: IBM-1 & IBM-2 Models

We reformulate the translation model with alignment:

$$Pr(f_1^J | e_1^I) = \sum_{a_1^J} Pr(f_1^J, a_1^J | e_1^I)$$

Further decomposed into a length model, alignment model and a lexicon model.

$$\begin{aligned} Pr(f_1^J, a_1^J | e_1^I) &= Pr(J | e_1^I) \cdot Pr(f_1^J, a_1^J | J, e_1^I) \\ &= Pr(J | e_1^I) \cdot Pr(a_1^J | J, e_1^I) \cdot Pr(f_1^J | a_1^J, J, e_1^I) \end{aligned}$$

Assumptions for IBM1 and IBM2 models:

- length model dependent only on length of target sentence:

$$Pr(J | e_1^I) = P(J | I)$$

- alignment model dependent only on absolute position and sentence length:

$$Pr(a_1^J | J, e_1^I) = \prod_{j=1}^J p(a_j | j, I, J)$$

- lexicon model dependent on aligned words only:

$$Pr(f_1^J | a_1^J, J, e_1^I) = p(f_j | e_{a_j})$$

Difference by alignment model:

- IBM-1 model: uniform probabilities:  $p(i | j, I, J) = \frac{1}{I+1}$
- IBM-2 model:  $p(i | j, I, J)$  is a large table which can be initialized by IBM-1 model

IBM-1 model consider all possible reordering as the same possibility. More often than not that, the word translation for input source word also follows the translation of the predecessor of that word, or more generally in limited distance. In IBM-2 model, we add an explicit model for alignment.



**First-Order Model: HMM**

Assume  $a_j$  depends on immediate predecessor position  $a_{j-1}$  as in speech recognition:

$$Pr(a_1^J | J, e_1^J) = \prod_j^J p(a_j | a_{j-1}, I, J)$$

**Weighted Model**

To overcome shortcomings in modeling, introduce scaling exponents  $\lambda_m$  as in speech recognition:

$$Q(f_1^J, e_1^I; a_1^J) = P(J|I)^{\lambda_1} \cdot \prod_i P(e_i | e_h)^{\lambda_2} \cdot \prod_j [p(a_j | a_{j-1}, I, J)^{\lambda_3} \cdot p(f_j | e_{a_j})^{\lambda_4}]$$

where  $e_h$  denotes the history words in  $N$ -gram model

**Fertility: IBM-3 Model**

Mostly, one word in the source language corresponding to just one single word in the target sentence. But fertility models that each source word can be translated to zero or more than one words.

**Fertility** models  $\varphi(e)$  the specific number of target words given an source word  $e$ . As consequence, length  $J$  depends on fertilities:

$$J = \sum_{i=0}^I \varphi(e_i)$$

Since fertility models the one-to-many translation or dropping input words. Similarly, we need to model adding words. IBM models these added words are generated by the special NULL token. We could model the fertility of the NULL token in the same way. After fertility step, we insert one NULL token with a specific probability  $p$  after each generated word.

IBM-3 model contains four steps:

- fertility step
- NULL insertion
- lexical translation step
- distortion step for reordering

### 2.1.2 Phrase-Based Model

Actually when translating, words may not be the best candidates for the smallest units for translation, sometimes one word in a foreign languages should be translated into two English words, or vice versa. Word-based models often break down in these cases.

Phrase-based models typically do not strictly follow the noisy-channel approach proposed for word-based models, but use a log-linear framework This model allow s straightforward integration if additional features.

**Log-linear Model Combination:** Consider arbitrary models ("feature functions"):

$$q_m(f_1^J, e_1^I; a_1^J) > 0 \quad m = 1, \dots, M$$

$$\begin{aligned} Q(e_1^I, f_1^J; a_1^J) &= \prod_{m=1}^M q_m(f_1^J, e_1^I; a_1^J)^{\lambda_m} \\ &= \exp\left(\sum_{m=1}^M \lambda_m \log q_m(f_1^J, e_1^I; a_1^J)\right) \end{aligned}$$

In this frame work, we view each data point as a vector of features and the model as a set of corresponding feature functions, these functions are trained separately and combined assuming that they are independent of each other.

Components for log-linear model can be such as language model, phrase translation model, reordering model are used as feature functions with appropriate weights.

- Phrase translation model can be learned from a word-aligned parallel corpus, alternatively we also can use expectation maximization algorithm to directly find phrase alignment for sentence pairs.
- Reordering model for phrase case is typically modeled by a distance based reordering cost that discourage reordering in general. Lexicalized reordering model can be introduced for specific phrase pair.

The model can be further extended by components like: bidirectional translation probabilities, lexical weighting, word penalty and phrase penalty.

## 2.2 Neural Machine Translation

Unlike traditional phrase-based machine translation, which consists of several models that are tuned separately, neural machine translation tries to build a more general neural network model which can directly output translations given input, it

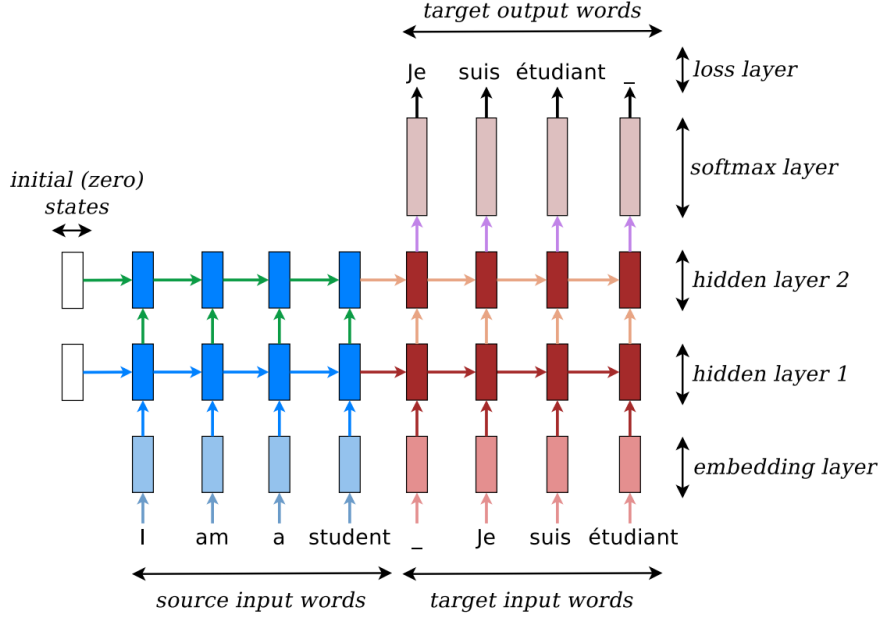


Figure 2.1: Neural machine translation – example of a deep recurrent architecture (Luong et al. [2015])

contains only one single model, and only one single training criterion. The first successful neural machine translation, attention mechanism has lately been used to improve neural machine translation by selectively focusing on parts of the source sentence during translation. The inherently sequential nature precludes parallelization within training examples.

### 2.2.1 Seq2seq Model

The most common network structure is the encoder-decoder framework, in which two recurrent neural networks work together to transform one sequence to another. An encoder condenses an input sequence into a vector, and a decoder unfolds that vector into a new sequence. The model can be expressed as:

$$p(e_1^I | f_1^J) = \prod_t p(e_i | e_0^{i-1}, f_1^J) = \prod_i p(e_i | e_0^{i-1}, \mathbf{s})$$

Since the decoder is actually a RNN structure, the conditional probability can be simplified as:

$$p(e_i | e_0^{i-1}, \mathbf{s}) = p(e_i | e_{i-1}, \mathbf{h}_{i-1}, \mathbf{s})$$

where  $\mathbf{h}_{i-1}$  denotes the hidden state of decoder and  $\mathbf{s}$  is the dense vector representation of source sentence. In more detail, the probability of decoding each word  $e_i$  as:

$$p(e_i|e_{i-1}, \mathbf{h}_{i-1}, \mathbf{s}) = \text{softmax}(g(\mathbf{h}_i))$$

$g$  is trainable function which gives the probability of all target words in the vocabulary.  $\mathbf{h}_i$  is the RNN hidden unit, abstractly computed as:

$$\mathbf{h}_i = f(\mathbf{h}_{i-1}, \mathbf{s})$$

$f$  defines the transformation connecting hidden states.

Drawbacks of the such seq2seq model:

1. The model compresses all information from the input sentence into a hidden vector 1, while ignores the length of input sentence, when the length of input sentence get very long, even longer than the training sentences, it becomes harder to extract specific information for predicting the target word, the performance will get worse.
2. It's not suitable to assign the same weight to all input words, one target word corresponds usually to one or several words in the input sentence. Treating all words equally does not distinguish the source information and influence the performance badly.

### 2.2.2 Attention Mechanism

To solve the problem mentioned above, attention mechanism was proposed to derive a context vector  $\mathbf{c}_i$  that capture the input information to help to predict the target word at the position  $i$ . The basic idea is: given the target hidden state  $\mathbf{h}_i$  and the source-side context vector  $\mathbf{c}_t$ , we can compute the hidden state  $\tilde{\mathbf{h}}_i$  by combining the current hidden state  $\mathbf{h}_i$  and the context vector  $\mathbf{c}_i$ :

$$\tilde{\mathbf{h}}_i = \tanh(W_c[\mathbf{c}_i; \mathbf{h}_i])$$

Then the target word can be predicted by softmax function:

$$p(e_i|e_0^{i-1}, f_1^J) = \text{softmax}(W_s \tilde{\mathbf{h}}_t)$$

The concept of attention mechanism comes first from the computer vision domain, when generating image captions, The model learns  $f$  to restrict attention to particular objects in the image. We denote the decoder RNN state as  $\mathbf{s}_i$ , the encoder hidden state over all position as  $\mathbf{h}_j$ .

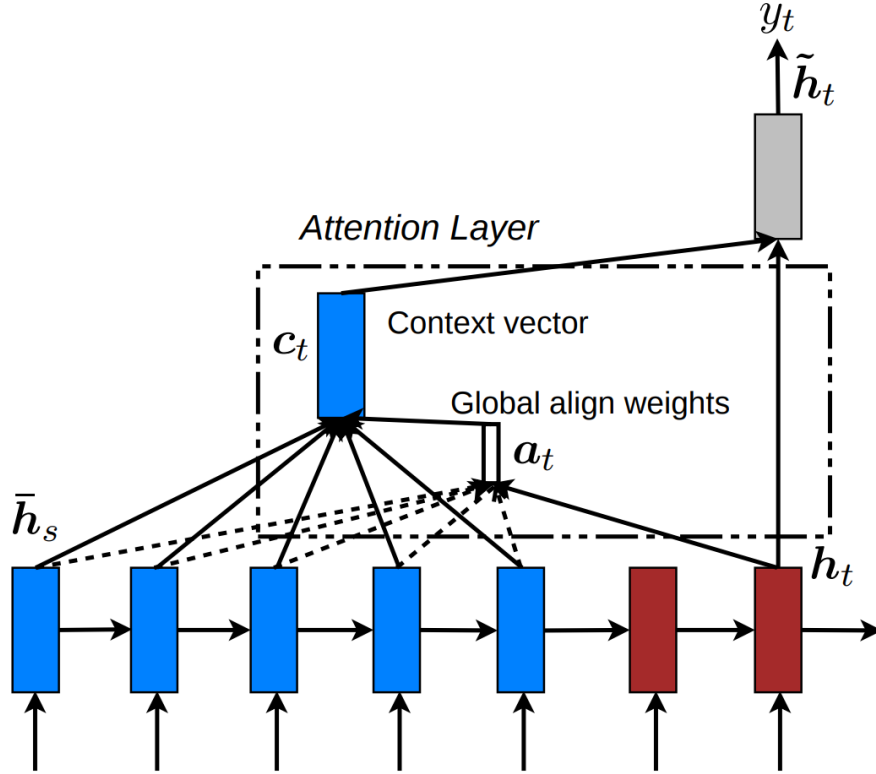


Figure 2.2: Global attention model (Luong et al. [2015])

### Global attention

The global attention attend all the input words, weighted average of source hidden states (word representations):

$$c_i = \sum_j \alpha_i(j) \cdot h_j$$

where  $\alpha_i(j)$  is the normalized attention weight that output at position  $i$  is aligned with input position  $j$ , and it can be calculated as the following softmax like function:

$$\alpha_i(j) = \text{align}(s_i, h_j) \tag{2.1}$$

$$= \frac{\exp(\text{score}_{\text{Xu et al. [2015]}}(s_i, h_j))}{\sum_{j'} \exp(\text{score}(s_i, h'_{j'}))} \tag{2.2}$$

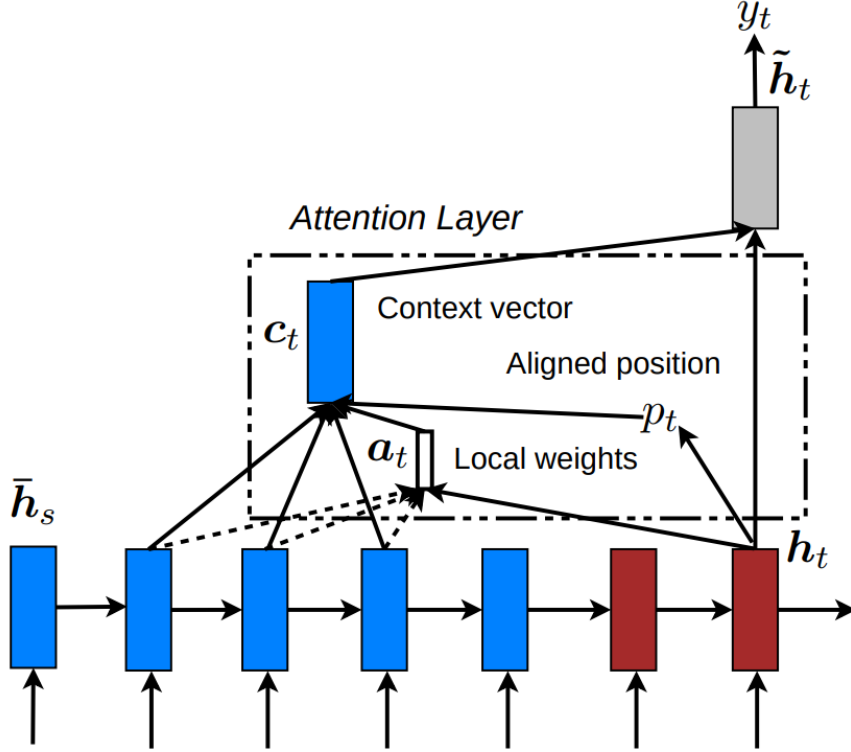


Figure 2.3: Local attention model (Luong et al. [2015])

Since score function is referred as a content based function, different forms can be considered:

$$\text{score}(\mathbf{s}_i, \mathbf{h}_j) = \begin{cases} \mathbf{s}_i^T \mathbf{h}_j & \text{dot} \\ \mathbf{s}_i^T W_a \mathbf{h}_j & \text{general} \\ \mathbf{v}_i^T \tanh(W_a[\mathbf{s}_i; \mathbf{h}_j]) & \text{concat} \end{cases} \quad (2.3)$$

### Soft attention & hard attention

With soft attention you need to calculate the attention over all features. While for hard attention, it is a deterministic method, instead of a weighted average, it uses  $\alpha_i$  as a sample rate to pick one feature. That's to say, in soft attention case: when computing the distribution of the alignment, for each word in the input sentence, the model will give a probability. Hard attention will determine a specific word from the input sentence as the alignment and force the alignment probability as 0. Hard attention mechanism works in image processing however it performs worse in text processing. Because such one-to-one alignment will produce bad translation

once a mis-alignment occurs.

### Local attention

Since for global attention, for each target word we need to attend the whole input sentence, it is very expensive and impractical to translate longer sentences. Luong et al. [2015]) proposed the local attention, the local attention is actually the trade-off between soft and hard attention.

For local attention, the model first generates an aligned position  $a(i)$  for each target word at position  $i$ . Then the context vector  $\mathbf{c}_i$  is a weighted sum within the window  $[a(i) - D, a(i) + D]$ ,  $D$  is selected empirically. The model predict the aligned position  $p_i$  as followed:

$$a(i) = S \cdot \text{sigmoid}(\mathbf{v}_p^T \tanh(W_p \mathbf{h}_i))$$

$W_p$  and  $\mathbf{v}_p$  are the model parameters which will be learned to predict positions.  $S$  is the source sentence length. As a result of sigmoid function,  $a(i) \in [0, S]$ . To favor the words near position  $a(i)$ , they place a Gaussian distribution which centered at  $a(i)$ .

$$\alpha_i(j) = \text{align}(\mathbf{s}_i, \mathbf{h}_j) \cdot \exp\left(-\frac{(j - a(i))^2}{2\sigma^2}\right)$$

### 2.2.3 Transformer

The RNN encoder-decoder models have achieved state of the art in sequence modeling and machine translation problem. However such RNN models also have some disadvantages, because of their inherently sequential computation which prevents parallelization across elements of the input sequence. That means it is more difficult to fully take advantage of the modern computing devices like GPU and TPU. Convolutional Gehring et al. [2017] and fully-attentional feed-forward architectures like Transformer Vaswani et al. [2017] models are proposed as alternatives for RNNs.

**Dot-Product Attention** Given a query  $\mathbf{q}$  and a set of key-value  $(\mathbf{k} - \mathbf{v})$  pairs, the output is weighted average of values, where the weight of each value is computed by inner product of query and corresponding key. The queries and keys are of the same dimension  $d_k$ , values are of dimension  $d_v$ . Actually, for better understanding, the query  $\mathbf{q}$  can be considered as the decoder state  $\mathbf{s}_i$  and key  $\mathbf{k}$  and values  $\mathbf{v}$  can be considered as the encoder state  $\mathbf{h}_j$  in previous attention model. Then the attention of query  $\mathbf{q}$  on all  $(\mathbf{k} - \mathbf{v})$  pairs are:

$$A(\mathbf{q}, K, V) = \sum_i \frac{\exp(\mathbf{q} \cdot \mathbf{k}_i)}{\sum_j \exp(\mathbf{q} \cdot \mathbf{k}_j)} \cdot \mathbf{v}_i$$

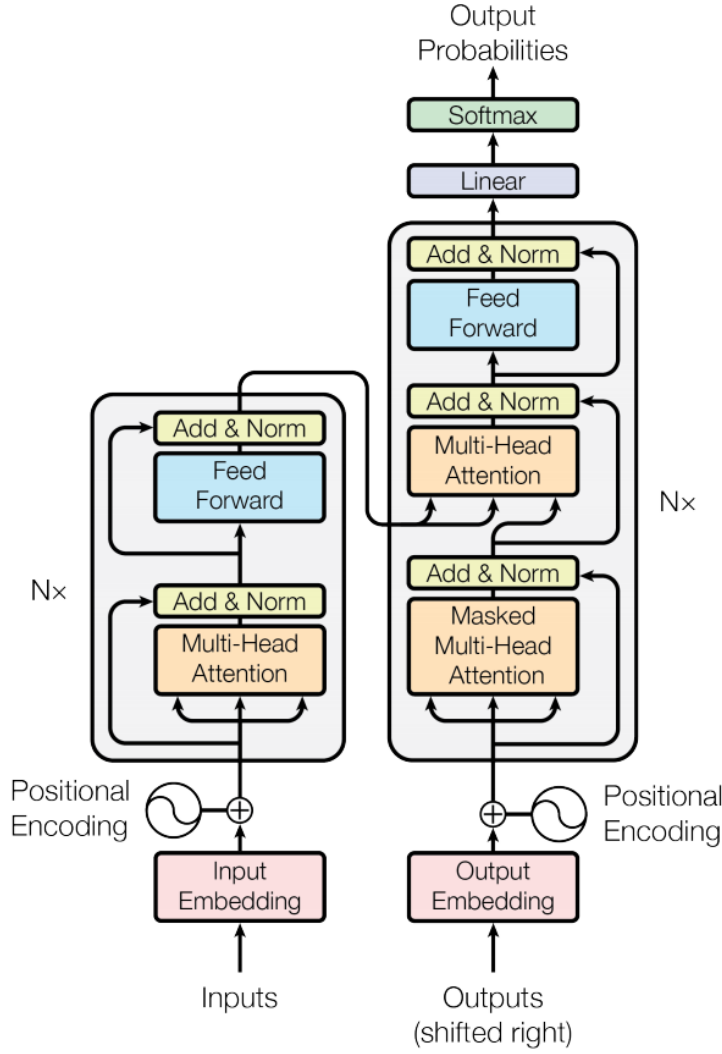


Figure 2.4: The Transformer model architecture (Vaswani et al. [2017])

When we stack the queries  $\mathbf{q}$  to  $\mathbf{Q}$ :

$$A(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}(\mathbf{Q}\mathbf{K}^T)\mathbf{V}$$

As  $d_k$  get larger, the variance of  $q^T k$  get larger. The softmax become very peaked and the gradient get smaller. To counteract this effect, we scaled the dot products by  $\frac{1}{\sqrt{d_k}}$ , we get



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

The shape of resulting matrix is  $n_q \times d_v$

### Multi-head attention

First map  $Q, K, V$  into  $h$  many lower dimension spaces via linear mapping matrix. Here  $h$  is the number of heads. Then apply attention and concatenate the outputs. Suppose the original dimensions of queries, keys, values are  $d_m$ . We the number of heads is  $h$ , then we set  $d_k = d_v = d_m/h$ . With  $i \in 1 \dots h$  different matrix  $W_i^Q \in R^{d_m \times d_k}$   $W_i^K \in R^{d_m \times d_k}$   $W_i^V \in R^{d_m \times d_k}$  and  $W^O \in R^{d_m \times d_k}$ .

$$\text{MultiHead}(Q, K, V) = \text{Concat}(h\text{head}_1, \dots, \text{head}_h)W^O$$

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

where  $\text{head}_i \in R^{n_q \times d_k}$ ,  $\text{Multihead}(Q, K, V) \in R^{n_q \times d_m}$

There are three attention types in the transformer model:

- encoder-decoder attention  
the queries come from the previous decoder layer, the keys and values are just the output of decoder. This works as the attention mechanism in the seq2seq model, it allow the decoder to align all positions in the input sequence with different weights
- encoder self-attention  
all queries, keys, values come from the encoder layer, each position allows to attend to all positions before that position
- decoder self-attention  
similar to encoder attention, each position is allowed to attend to position before and including that position

### Positional Encoding

Since there is no RNN or CNN structures in transformer model, we need also need to make use of sequence information for seq2seq learning. We need inject position information into the embedding.

$$PE_{(\text{pos}, 2i)} = \sin(\text{pos}/10000^{2i/d_m})$$

$$PE_{(\text{pos}, 2i+1)} = \cos(\text{pos}/10000^{2i/d_m})$$

where  $\text{pos}$  is the position and  $i$  is the dimension. That is, each dimension of the positional encoding corresponding to a sinusoid.

## 2.3 Unsupervised Machine Translation

give the first results for training a full translation model from non-parallel data, and use this model to translate previously-unseen text. The proposed problem named "decipherment" has received a lot of attention. This model contains many statistical methods, like EM algorithm or Bayesian method. However, this method is limited to rather short sentences and simplistic setting. With the development of neural machine translation, the end-to-end method performs better and easier to tune without much specific linguistic knowledge. The Idea is similar to dual learning framework except that in dual learning model, gradients are backpropagated through the reverse model and pretrain using a relative large parallel data as a warm start. The neural model mapping the sentences from monolingual corpora into the same latent space, By learning to reconstruct in both languages from the shared feature space. The translation model in both direction will be improved synchronously with the back-translation.

### 2.3.1 Decipherment

For simplification, we use  $\mathbf{e}, \mathbf{f}$  instead of  $e_1^M, f_1^N$  to describe the sentence pairs. IBM model tries to maximize the probability with hidden alignment model

$$\operatorname{argmax}_{\theta} \prod_{\mathbf{e}, \mathbf{f}} P_{\theta}(\mathbf{f}|\mathbf{e}) = \operatorname{argmax}_{\theta} \prod_{\mathbf{e}, \mathbf{f}} \sum_a P_{\theta}(\mathbf{f}, a|\mathbf{e})$$

While for unsupervised case, we train

$$\operatorname{argmax}_{\theta} \prod_{\mathbf{f}} P_{\theta}(\mathbf{f}) = \operatorname{argmax}_{\theta} \prod_{\mathbf{f}} \sum_{\mathbf{e}} P(\mathbf{e}) \cdot P_{\theta}(\mathbf{f}|\mathbf{e})$$

for hidden alignments:

$$\operatorname{argmax}_{\theta} \prod_{\mathbf{f}} \sum_{\mathbf{e}} P(\mathbf{e}) \cdot \sum_a P_{\theta}(\mathbf{e}, a|\mathbf{e})$$

Since the model is very complicated, more assumptions are added to the model. The model accounts for word substitutions, insertions, deletions and local reordering during the translation process but does not incorporate fertilities or global re-ordering.

The generative process:

1. Generate an English sentence  $\mathbf{e}$  with probability  $P(\mathbf{e})$ .
2. Insert a NULL word at any position in the English sentence with uniform probability.

3. For each English word token  $e_i$  (including NULLs), choose a foreign word translation  $f_i$ , with probability  $P_\theta(f_i|e_i)$ , the foreign word may be NULL.
4. Swap any pair of adjacent foreign words  $f_{i-1}, f_i$ , with probability  $P_\theta(\text{swap})$ .
5. Output the foreign string  $f_1^M$ , skipping over NULLs.

### 2.3.2 Neural Unsupervised Machine Translation




Figure 2.5: Illustration of unsupervised machine translation: A) the two original monolingual data; B) Initialization, the two distribution are roughly aligned; C) Denoising autoencoder, make the distribution closer to normal one in corresponding language D) Back-translation, improve the both translation model in iterative way (Lample et al. [2018])

#### Dual Structure

Dual structure is inspired by the observation: any machine translation has a dual task: e.g. German-to-English translation (primal) and English-to-German translation (dual). Dual tasks can form a closed loop and generate informative feedback signals to for both translation models. Based on the feedback signals generated during this process, and leverage language model, we can minimize the reconstruction error of the original sentences. We can iteratively update the two models until convergence. He et al. [2016] trained two agents to translate in opposite directions (e.g. French  $\rightarrow$  English and English  $\rightarrow$  French), and make them teach each other through a reinforcement learning process. While promising, this approach still requires a parallel corpus for a warm start. Lample et al. [2018] simplified the structure, in the proposed model, the gradients will not be back-propagated through the reverse model, but only make use of the back-translation. The training process as followed:

---

**Algorithm 1** Unsupervised Machine Translation

---

**Language models:** Learn language models  $P_s, P_t$  over source and target languages;

**Initial translation model:** Leveraging  $P_s, P_t$ , learn two initial translation models in each direction:  $P_{s \rightarrow t}^{(0)}, P_{t \rightarrow s}^{(0)}$ ;

**for**  $k = 1$  **to**  $N$  **do**

**Back-translation:** Generate source and target sentences using the current translation models  $P_{t \rightarrow s}^{(k-1)}$  and  $P_{s \rightarrow t}^{(k-1)}$ , leveraging  $P_s, P_t$ ;

**Retrain:** Train new translation models  $P_{t \rightarrow s}^{(k)}$  and  $P_{s \rightarrow t}^{(k)}$  using the generated sentences and leveraging  $P_s, P_t$ ;

**end**

---

**Initialization**

If without enough information to start the dual machine translation system, it will be hard for the system to catch the meaningful signals and then it will take much more iterations. However if we initialize the system by a naive a word-by-word translation of the sentence, where the bilingual lexicon are derived from the same monolingual data. While such initial "word-by-word" translation maybe poor if languages or corpora are not closely related, it still preserves some of the original semantics. Such word-by-word translation can already achieve several BLEU scores.

**Shared Latent Representation**

A shared encoder representation acts like an interlingua, which is translated in the decoder corresponding language regardless less of the input source language. Since the only supervision information only comes from the monolingual data. The model learn to translate by learning to reconstruct in both language from this shared space.

There are two different methods to learn the shared feature space:

- **Shared encoder** The system use only one encoder that is shared by both languages involved. The universal encoder is aimed to produce a language independent representation of the input language but the decoder should separately translate then into corresponding language.
- **Adversarial training** Train discriminator to classify between the encoding of source and target sentences. The discriminator operates on the output of the encoder, the encoder is trained instead to fool the discriminator.

**Optimization**

When minimizing the loss function, gradients will not be back propagated through the reverse model which generate the data. Instead the objective function minimized at every iteration is the sum of  $L^{auto}$  and  $L^{back}$

- Denoising autoencoder loss

$$L^{auto} = \mathbb{E}_{e \sim E}[-\log P_{t \rightarrow t}(e|\text{noise}(e))] + \mathbb{E}_{f \sim F}[-\log P_{s \rightarrow s}(f|\text{noise}(f))]$$

where  $s$   $P_{s \rightarrow s}$  and  $P_{t \rightarrow t}$  are the composition of encoder and decoder both operating in the source and target sides, respectively

- Back-translation loss

$$L^{back} = \mathbb{E}_{e \sim T}[-\log P_{s \rightarrow t}(e|u(e))] + \mathbb{E}_{f \sim S}[-\log P_{t \rightarrow s}(f|v(f))]$$

we denote the sentence that translated by intermediate target-to-source translation model as  $u(y)$ , similarly denote the sentence translated by source-to-target model as  $v(x)$ , so  $u(y)$  should in source language and  $v(x)$  should in target language. The pairs  $(x, v(x))$ ,  $(u(y), y)$  constitute synthetic parallel sentences.



## Chapter 3

# Word Embedding

Word embeddings is distributed representation of words in a vector space. With the learning algorithm it can capture the contextual or co-occurrence information. The word embedding has an interesting and important property: similar words will have similar distribution in the embedding space, with that property, we can find meaningful near-synonyms or Some successful methods for learning word embedding like word2vecMikolov et al. [2013c], Pennington et al. [2014]

### 3.1 Monolingual Embedding

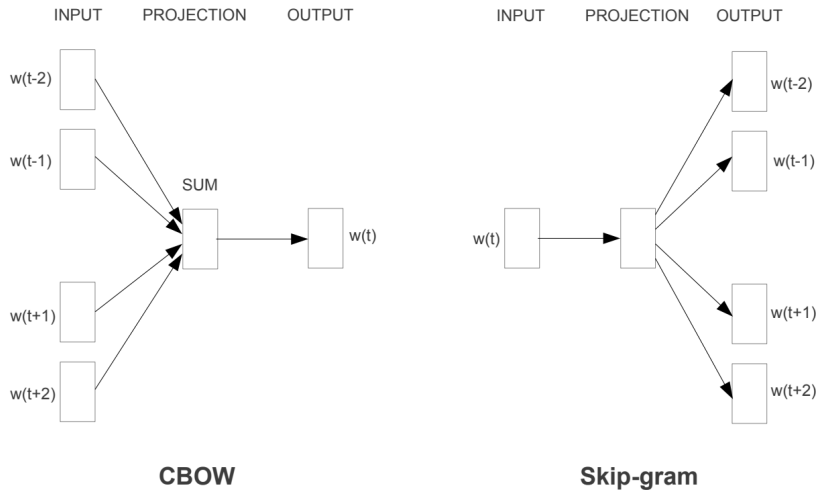


Figure 3.1: Global attention model (Mikolov et al. [2013a])

#### 3.1.1 CBOW and Skip-gram Model

CBOW model and skip-gram model are currently mainstream neural network to learn the word embedding. Algorithmically, CBOW tries to predict the current

word based on the context while skip-gram model tries to maximize classification of a word based on another word in the same sentence. Using skip-gram model as example, given a large training corpus represented as a sequence of words  $w_1, \dots, w_N$ , the objective of the model is to maximize the following log-likelihood:

$$\sum_{s=1}^N \sum_{w_t \in \mathcal{C}_s} \log p(w_t|w_s)$$

where  $w_s$  is the current word, the context  $\mathcal{C}^s$  is the set of words surrounding word  $w_s$ , so  $w_t$  means the word to predict. So probability above can be interpreted as observing a context word  $w_t$  given  $w_s$ .

Suppose we can given a scoring function  $s$  which maps word pairs to score in  $\mathbb{R}$  like cosine similarity, the probability can be defined like softmax function:

$$p(w_t|w_s) = \text{softmax}(s(w_t, w_s)) \quad (3.1)$$

$$= \frac{\exp\{s(w_t, w_s)\}}{\sum_{w' \in W} \exp\{s(w', w_s)\}} \quad (3.2)$$

where  $W$  is the whole vocabulary.

We train the model by maximizing its log-likelihood on the training dataset, i.e. by maximizing:

$$J_{ML} = \log p(w_t|w_s) \quad (3.3)$$

$$= s(w_t, w_s) - \log\left(\sum_{w' \in W} \exp\{s(w', w_s)\}\right) \quad (3.4)$$

### Noise-Contrastive Training

However the normalization on the whole vocabulary is very expensive since it is conducted for all words at every training step. The problem of predicting words can be considered as an independent binary classification task. For example in the skip-gram model, we consider all the context words as positive samples and the words randomly sampled from the dictionary as the negative ones. Then the training objective is

$$J_{NEG} = \log Q_\theta(D = 1|w_t, w_s) + \sum_{w' \sim P_{noise}} \log Q_\theta(D = 0|w', w_s)$$

where  $Q_\theta(D = 1|w_t, w_s)$  is the binary logistic regression probability. In practice, we draw  $k$  contrastive words from the noise distribution. Then the model are instead trained to discriminate the read target word  $w_t$ , from imaginary (noise) words.



Since we only calculate the loss function for  $k$  samples instead the whole vocabulary, it becomes much faster to train.

According to empirical results, CBOW works better on smaller datasets because CBOW smoothes over a lot of the distributional information while Skip-Gram model performs better when we have larger datasets

### 3.1.2 FastText

The training methods above treat each word as a distinct word embedding, however intuitively we can obtain more information from the morphological information of words. A subword model was proposed to try to fix such problem. The training network is similar, the model design a new presentation of the word: it adds special symbols  $<$ ,  $>$  as boundary information at the beginning and the end of a word. Then a normal word is represented as a bag of character  $n$ -grams. For example the word "where" and  $n$  equals 3, then it can be represented as the following 5 tri-grams:

$$< wh, whe, her, ere, re >$$

Suppose in this way we denote a word  $w$  as  $G_w$  the set of character  $n$ -grams, we assign for each character  $n$ -gram  $g$  in  $G_w$ , we assign a distinct vector  $z_g$ , we will finally represent the embedding of word  $w$  as the sum of these vector and also for the scoring function:

$$s(w, w_s) = \sum_{g \in G_w} z_g^T w_s$$

## 3.2 Cross-lingual Word Embedding

Cross-lingual word embedding is defined as word embedding of multiple languages in a joint embedding space. Mikolov first notice that the embedding distributions exhibit similar structure across languages. They proposed to use a linear mapping from the source embedding to target embedding.

Cross-lingual word embeddings are appealing due to two main factors: they enable not only cross-lingual semantics, reasoning about word meaning in multilingual contexts, for example, we can induce the bilingual dictionary by calculating the cross-lingual word similarities, but also knowledge transfer between languages, most importantly between rich-resource and lean-resource languages (Adams et al. [2017]).

In the thesis, I assume there are two set of embeddings  $e, f$  trained separately on monolingual data. The propose of cross-lingual word embedding training is to learn such a mapping  $W \in$  from source embedding space to target embedding space, so




Figure 3.2: A cross-lingual embedding space between German and English (Ruder et al. [2017])

$W_{f_i, e_j}$  in the same embedding space and for all corresponding word pairs, we need to optimize the mapping  $W$ , so that”

$$\arg \min_{W \in R^{d \times d}} \sum_i \|W f_i - e_i\|$$

where  $d$  is the dimension of embeddings, and the distance  $\|Wf_i - e_i\|$  can be different types. We prefer the Euclidean distance.

first observe that word embeddings trained separately on monolingual corpora exhibits isomorphic structure across languages, as illustrated in Figure . That means we can create a connection between source embedding and target embedding even with simple linear mapping. This has far-reaching implication on low-resource scenarios , because word embedding requires only plain text to train, which is the most abundant form of linguistic resource.

### 3.2.1 \*\*\*

According to the training method we can divide the supervised method into three: Multilingual learning can be categorized into mapping-based approaches and regularization-based approaches. In the mapping-based approaches, the embedding is performed for each language individually with monolingual data, and then the mapping are learned using multilingual data to represent the relation between the languages.

One method is to learn the linear mapping projection, another is to learn mappings that project word vectors of all languages to a common low-dimension space, where the correlation of the multilingual word pairs is maximized with the canonical correlation analysis (CCA). The advantages of approach is that it is very fast to learn the embedding alignments. The main draw back of this approach is that it is not clear if a single transformation whether linear or nonlinear can capture the relationship between all words in different languages.

The regularization-based approaches involve the multilingual constraint in the objective function for learning the embedding, adds an extra term that reflects the distances of some pairs of semantically related words from different languages into the objective function. Max-margin hinge loss (MMHL) learn word embeddings by train a model to output a higher score for a correct word sequence than for an incorrect one. Jointly optimize the monolingual objective  $\mathbb{L}(\cdot)$ , with the cross-lingual objective enforced as a cross-lingual regularizer. need to define a cross-lingual regularization term  $\Omega(\cdot)$  to constrain monolingual models as they are jointly being trained over the context

The joint formulation of the learning objective provided as:

$$J = \mathcal{L}^s + \mathcal{L}^t + \Omega(s, t)$$

where  $\mathbb{L}^s$  and  $\mathbb{L}^t$  are monolingual objectives in each language optimized jointly with the cross-lingual regularization objective. A dictionary is necessary for learning the cross-lingual word embedding. minimizing the distance in a bilingual dictionary.

Xing showed that the results are improved when we constrain the  $W$  to be an orthogonal matrix. This constraint, the optimal transformation can be efficiently calculated in linear time with respect to the vocabulary size.

The problem then is simplified as the Procrustes problem and there exists a closed-form solution obtained from the SVD of  $EF^T$

### 3.2.2 Orthogonal Constraints

Starting from Mikolov et al. [2013b], the mapping from source embedding space to target embedding space can be represented as a linear repression. The objective can be defined as:

$$\min_W \sum_i \|Wf - e\|^2$$

Since we retrieve the word translation according to cosine similarity, it's better to solve the problem by redefine the optimization function using the cosine distance:

$$\operatorname{argmax}_W \sum_i (Wf_i)^T e^i$$

. We consider the source and target embedding in the same space. In this case, the normalization constraint on word vectors can be satisfied by constraining  $W$  as an orthogonal matrix. This is equivalent to minimizing the (squared) Frobenius norm of the residual matrix:

$$W^* = \underset{W}{\operatorname{argmin}} \|WF - E\|_F^2$$

The problem boils down to the Procrustes problem which has a closed form solution obtained from the singular value decomposition (SVD).

$$W^* = UV^T, \quad U\Sigma V^T = \operatorname{SVD}(EF^T)$$

### 3.2.3 CSLS Loss

Inspired from the work of Conneau et al. [2017], where the dictionary induced from CSLS loss:  $e$

$$\operatorname{CSLS}(e, f) = -2\cos(e, f) + \frac{1}{k} \sum_{e' \in N_e(f)} \cos(e', f) + \frac{1}{k} \sum_{f' \in N_f(e)} \cos(f', e)$$

since we have  $\cos(\mathbf{W}e, f) = e^T \mathbf{W}^T f$

The loss function can be rewritten as:

$$\min_{\mathbf{W} \in} = \frac{1}{n} \sum_{i=1}^n$$

Minimization of a non-smooth cost function over the manifold of orthogonal matrices . Instead of using manifold optimization tools, ? proposed to derive convex relaxations that can lead to a simple and tractable minimization algorithm.

- Spectral norm: replacing the set of orthogonal matrices  $\mathbf{1}$  by its convex hull, that is the set of matrices with singular values smaller than 1, the unit ball of the spectral norm
- Frobenius norm: replacing the the set of orthogonal matrices  $\mathbf{1}$  with ball of radius  $\sqrt{d}$  in Frobenius norm

With such two relaxations, the CSLS loss is constrained to a convex function with respect to the mapping  $\mathbf{W}$ .

We train the linear mapping with in the spectral norm by projected gradient descent:2 For each iteration, train the mapping  $\mathbf{W}$  with gradient descent, then constrain the mapping by projection of the set.

- Spectral norm  
take the SVD of the trained matrix, threshold the singular values to one
- Probenius norm  
divide the matrix by its Frobenius norm

## Chapter 4

# Cross-lingual Word Embedding without Parallel Data

Proposed cross-lingual embedding learning requires several thousands of word pairs as anchors to learn the mapping with good generalization ability to predict for unseen word. But as for low-source language, dictionary of good quality is not readily available. This motivates more and more works on unsupervised cross-lingual word embedding learning, if the unsupervised method works, the word-level or grammar-level properties will be better captured, since it enables the knowledge transfer for those less-studied languages.

### 4.1 Initialization

Vulic and Korhonen [2016] shows that a seed dictionary of at least hundreds are need for the model to generalize. Hoshen and Wolf [2018] proposed an approach which first use PCA to realize an approximate alignment, then use mini-batch cycle iterative closest point to learn the mapping. Zhang et al. [2017] proposed an adversarial autoencoder, using the generator  $G$  to implement the mapping so that the transformed source embedding similar to the corresponding target embedding, and the discriminator  $D$  strives to distinguish the target embedding from the transformed source embedding. Conneau et al. [2017] simplifies the structure and proposed an unsupervised criterion that is highly correlated with the quality of the mapping which can be used as stop criterion and to select the best hyper-parameters.

#### 4.1.1 Iterative Closest Point Method

The method assumes that many language pairs share some principle axes of variation. For each language, first select most frequent word vector, center the data and project it to the top  $p$  principle components. The normal Iterative Closest Point (ICP) is unidirectional, the Mini-batch Cycle ICP (MBC-ICP) includes cycle-constraints ensuring that a word transformed to another language could also

be transformed back and stay unchanged. Each iteration the final MBC-ICP algorithm can be described as:

1. For each target word  $e$ , find the nearest  $W_{s \rightarrow t}f(e)$ , denote the source word  $f(e)$
2. For each source word  $f$ , find the nearest  $W_{t \rightarrow s}e(f)$ , denote the target word  $e(f)$
3. Optimize  $W_{e \rightarrow f}$  and  $W_{f \rightarrow e}$ :

$$\begin{aligned} \mathcal{L} = & \sum_e \|e - W_{s \rightarrow t}f(e)\| + \sum_f \|f - W_{t \rightarrow s}e(f)\| \\ & + \lambda \sum_e \|e - W_{s \rightarrow t}W_{t \rightarrow s}e\| + \lambda \sum_f \|f - W_{t \rightarrow s}W_{s \rightarrow t}f\| \end{aligned}$$

#### 4.1.2 Adversarial training

Generative adversarial networks are a class of artificial intelligence algorithm used in unsupervised learning. The objective of generative network (generator) is to increase the error rate of the discriminative network (discriminator) while the discriminator discriminates between instances from the true data distribution and candidates produced by the generator. Following this principle, the generator tries to learn to mapping from the source embedding space to target embedding space, so that the discriminator while could be a deep neural network cannot distinguish the data source.

Let  $= \{x_1, \dots, x_n\}$  and  $= \{y_1, \dots, y_m\}$  be the two sets of  $n$  and  $m$  word embeddings from a source and a target language separately, We refer the discriminator parameters as  $\theta_D$ . The discriminator is a multi-layer neural network trained to discriminate the transformed source word embedding from the target word embedding, while the mapping  $W$ , simply a linear transformation, is trained to fooling discriminator. In the two-player game, we are supposed to learn the mapping from source embedding space to the target space.

In the first adversarial training method for cross-lingual word embedding, the objective of the two networks are as followed:

Discriminator objective

$$L_D(\theta_D|W) = -\frac{1}{n} \sum_{i=1}^n \log P_{\theta_D}(\text{source} = 1|Wf_i) - \frac{1}{m} \sum_{i=1}^m \log P_{\theta_D}(\text{source} = 0|e_i)$$

Generator objective

$$L_W(W|\theta_D) = -\frac{1}{n} \sum_{i=1}^n \log P_{\theta_D}(\text{source} = 0|Wf_i)$$

To relax the orthogonal constraint, the authors introduce the adversarial autoencoder, after the mapping  $W$  from source embedding space to target embedding space, the source embedding should also be mapping back with  $W^T$ , therefore, they introduce the reconstruction loss as

$$L_W(W|\theta_D) = -\frac{1}{n} \sum_{i=1}^n \{\log P_{\theta_D}(\text{source} = 0|Wf_i) - \lambda \cos(f_i, W^T W f_i)\}$$

In Conneau et al. [2017] work, they just take a symmetric loss for the generator loss and achieve better results:

$$L_W(W|\theta_D) = -\frac{1}{n} \sum_{i=1}^n \log P_{\theta_D}(\text{source} = 0|Wf_i) - \frac{1}{m} \sum_{i=1}^m \log P_{\theta_D}(\text{source} = 1|e_i)$$

### Model Selection

Since the cross-lingual embedding training is under the unsupervised setting, We do not know the word translation accuracy, otherwise if we have the validation data, that means we will have parallel data, against the unsupervised idea. To address this issue, we must select from the property of data or the loss of the neural network as the unsupervised criterion. However in the experiments we find that the accuracy of the discriminator always stays at a high level no matter how is the word translation accuracy.

All these methods can be use to find meaningful word pairs in both languages. For further refinement we can use the induced dictionary to start the iterative self-learning algorithm.

## 4.2 Nearest Neighbor Search

### Hubness Problem

Points are tending to be nearest neighbors of many points in high-dimensional space. Since we use nearest neighbor search, those points (hubs) will harm the search accuracy.

### Inverted Softmax

The confidence of choosing a target word as translation of a source word can be considered as softmax-like normalized probability

$$p(e|f) = \frac{\exp(\beta \cdot s(e, f))}{\sum_{e'} \exp(\beta \cdot s(e', f))}$$

where the  $s(\cdot)$  is the score function we can define ourselves. we learn the "temperature"  $\beta$  by maximizing the log probability over the training dictionary.

$$\operatorname{argmax}_{\beta} \sum_{e,f} \ln p(e|f)$$

The author observed that, if we invert the softmax and normalizing the probability over all the source words rather than target words, the hubness problem could be mitigated:

$$p(e|f) = \frac{\exp(\beta \cdot s(e, f))}{\sum_{f'} \exp(\beta \cdot s(e, f'))}$$

### Cross-domain Similarity Local Scaling

We denote  $N(f)$  the set of  $K$  nearest neighbors of points of the mapped  $f$  in the target embedding space, and  $N(e)$  the nearest neighbors of mapped  $e$  in the source embedding space. We consider the mean cosine similarity as hub-ness:

$$r(f) = \frac{1}{K} \sum_{e' \in N(f)} \cos(e', Wf)$$

So in this way we penalize the hub points:

$$CSLS(e, f) = 2\cos(e, y) - \frac{1}{K} \sum_{e' \in N(e)} \cos(f, e') - \frac{1}{K} \sum_{f' \in N(f)} \cos(f', e)$$

With the nearest neighbor search we can evaluate the performance of learned mapping, by checking if the nearest neighbors of the source word are also in the candidates in a open word translation dictionary dataset.

## 4.3 Iterative Training

Assume that with the initialization step we can already get roughly aligned distribution of the source and target embedding, and we can use this mapping to learn a good dictionary (word pairs). Since the quality of dictionary gets improved, we can further learn a better mapping, consequently, such process can be repeated iteratively until convergence criterion is met.



**Algorithm 2** Iterative training procedure

---

**Input:**  $\mathcal{F}$  (source embeddings)  
**Input:**  $\mathcal{E}$  (target embeddings)  
**Input:**  $\mathcal{D}$  (seed dictionary)  
**Result:**  $\mathcal{W}$  (embedding mapping)  
**while** *not converge* **do**  
     $\mathcal{W} \leftarrow \text{LEARN\_MAPPING}(\mathcal{F}, \mathcal{E}, \mathcal{D})$   
     $\mathcal{D} \leftarrow \text{LEARN\_DICTIONARY}$   
**end**

---

**4.3.1 Self-learning with Procrustes**

For self-learning frame work, we compute the optimal dictionary based on the mapped source word embeddings and target word embeddings. To sure the quality of the small dictionary, we choose only top frequent words for both languages and only bidirectional translations are kept. We even filter the dictionary with some threshold.

**4.3.2 Data-driven Method**

Since we have large monolingual data and one of the tasks to evaluate the cross-lingual word embedding is build machhine translation system based on these cross-lingual embeddings. We make use of the language model and improve also the translation quality simultaneously. We can set the source from the translation, build the dictionary from the word translation pairs, in detail the training procedure are as followed:

1. translate corpus according to current mapping, get the word pairs  $D$
2. train the network with  $D$  to minimize the mapping distance
3. Repeat 1, 2 until the algorithm converges

$$\left. \begin{array}{c} (f_1, e_1) \\ (f_2, e_2) \\ \vdots \\ (f_N, e_N) \end{array} \right\} \Rightarrow D$$

Previously, the objective we want to minimize is the mean square error of the mapped embeddings. By training the mapping with neural network we can define different loss functions according to specific features

$$\mathcal{L} = \sum_{(f,e) \in D} \|Wf - e\|^2$$



# Chapter 5

## Sentence Translation

As mentioned previously, training a traditional machine translation system requires large parallel data. With cross-lingual word embedding we can already found ambiguous word translation, in this chapter we propose a simple yet effective method to improve quality of translation which starts from the word-by-word translation. We integrate monolingual model like language model and denoising neural network of the target side to produce meaningful sentence translation. Since all the sub-models are trained on monolingual corpora, our proposed method is total unsupervised. Our system surpasses state-of-the-art unsupervised translation without costly iteratively training.

### 5.1 Context-aware Beam Search

#### 5.1.1 Language Model

Language models are widely applied in natural language processing, especially machine translation which need frequent queries.

*n*-gram models

*N*-gram language models use the Markov assumption to break the probability of a sentence into the product of the probability of each word given a limit history of preceding words.

$$p(w_1^N) = \prod_{i=1}^N p(w_i | w_1, \dots, w_{i-1}) = \prod_{i=1}^N p(w_i | w_{i-(n-1)}, \dots, w_{i-1})$$

The conditional probability can be calculated from *n*-gram model frequent counts:

$$p(w_i | w_{i-(n-1)}, \dots, w_{i-1}) = \frac{\text{count}(w_{i-(n-1)}, \dots, w_i)}{\text{count}(w_{i-(n-1)}, \dots, w_{i-1})}$$

Language model tries to handling sparse data problem because some words or phrases have not been seen yet in the training corpus does not mean they are not impossible. Different smoothing techniques like back-off or interpolation are implemented to assign a probability mass to unseen cases.

### 5.1.2 Beam Search

The complexity of a search graph is exponential to the length of the given source sentence. Beam search is a heuristic search algorithm that explores a graph by expanding the most promising nodes. At each step of the search process, it will evaluate all the candidates together with the reserved translation results from last step, it will only store a predetermined number (beam width) of translations for next step. The greater the beam width is, the fewer states will be pruned. So it is suggested to prune these word translation candidates as soon as possible to reduce the search space and speed up the translation. According to the similarity of cross-lingual word embedding, we are able to find some meaningful for translation candidates for a given word. But there are also words that actually noise in the candidates or obviously incorrect because of grammar checking. With the support of language model, we can select the most probable words from previous word translation candidates.

Given a history  $h$  of target word before  $e$ , the score of  $e$  to be the translation of  $f$  is defined as:

$$\hat{e}_1^N = \operatorname{argmax}_{e_1^N} \prod_{n=1}^N p^{\lambda_{LM}}(e_n | e_{n-4}^{n-1}) \cdot q^{\lambda_{emb}}(f_n, e_n)$$

where the lexicon score  $q(f, e) \in [0, 1]$  defined as:

$$q(f, e) = \frac{d(f, e) + 1}{2}$$

$d(f, e) \in [-1, 1]$  cosine similarity between  $f$  and  $e$

In experiments, we find such lexicon score works better than others, e.g. sigmoid or softmax

## 5.2 Denoising Neural Network

### 5.2.1 Denoising Auto-encoder

With the help of language model we have actually improved the quality of word-by-word translation but the results are still far from a acceptable one because of the drawback of the word-by-word mechanism, maybe to some degree we can infer the meaning of sentence, but the sequence of sentences depends on the specific language. We implement the sequential denoising autoencoder to improve the translation output.

An autoencoder is a neural network that is trained to copy its input to its output, autoencoders minimize the loss function like:

$$L(\mathbf{x}, g(f(\mathbf{x})))$$

where  $L$  penalizing the difference between the input and output. While a denoising autoencoder (DAE) instead minimizes

$$L(\mathbf{x}, g(f(\tilde{\mathbf{x}})))$$

where  $\tilde{\mathbf{x}}$  is a noise transformation of  $\mathbf{x}$  and denoising autoencoder will try to learn to ignore the noise in  $\mathbf{x}$  reconstruct the correct one. Sequential denoising autoencoder will find robust representation of sentences. In practice, denoising autoencoder consists of two parts, namely encoder and decoder. The encoder processes noised data and produces real-valued vectors as an encoding feature of the data. The computational graph of the denosing autoencoder, which attempt to reconstruct the normal input  $\mathbf{x}$  from it corrupted version  $\tilde{\mathbf{x}}$ . The model is trained by minimize the loss.

For our sequential denoising model, the label sequences would be the monolingual data of the target language. However we do not have the noise input. In order to make the model run correctly, we should mimic the noise sentence of word-by-word translation on the target monolingual corpus.

We design different noise types w.r.t. the word-by-word translation. In the experiments, we inject the artificial noise into a clean sentence, the experiment results shows the noise is reasonable and suitable in this case.

### 5.2.2 Noise Model

We design three types of noise to handle the fertility and reordering problem, namely reordering noise, insertion noise and deletion noise. In experiments, the noise model can improve the sentence translation, but since it actually starts from the word-by-word translation, it can only deal with reordering in limited distance, cannot work for global reordering.

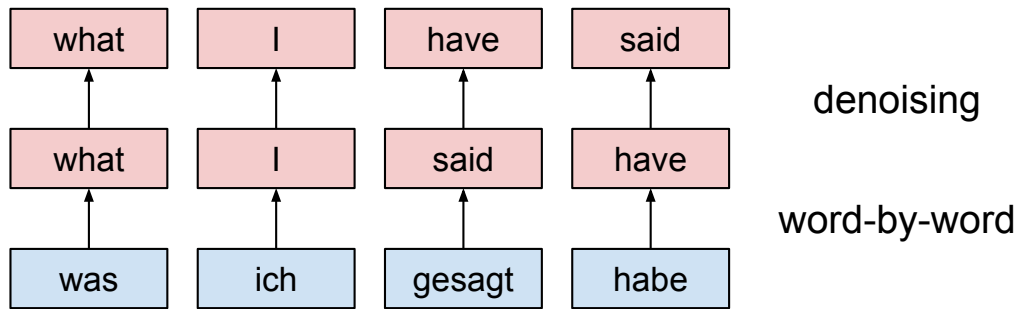


Figure 5.1: Reordering noise

## Reordering Noise

The reordering problem is a common phenomenon in the word-by-word translation since the sequence in source language is not exact the sequence in target language. For example, in the grammar of German, the verb is often placed at the end of the clause. "um etwas zu tun". However in English it is not the case, the corresponding translation sequence is "to do something". The verb should always before the noun. In our beam search, language model only assisting in choosing more suitable word from the translation candidates, it cannot reorder the word sequence at all.

For a clean sentence from the target monolingual corpora, we corrupt the word sequence by permutation operation. We limit the maximum distance between the original position and its new position.

The design of reordering noise is as followed:

1. For each position  $i$ , sample an integer  $\delta_i$  from  $[0, d_{per}]$
2. Add  $\delta_i$  to index  $i$  and sort  $i + \delta_i$
3. Rearrange the words to be in the new positions, to which where indices have been moved

Reordering is actually depends on the specific language pair. However in the experiments we found the performance of the denoising network aimed at such noise is not obvious. The Bleu score before and after the process is close.

## Insertion Noise

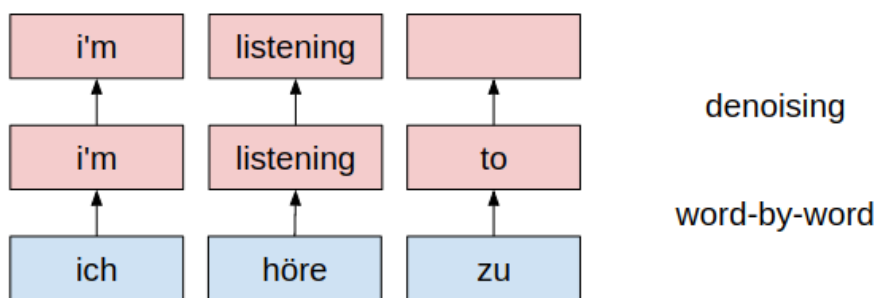


Figure 5.2: Insertion noise

The word-by-word translation system predict the source word at every position of the sentence. However the vocabularies of different systems are not symmetric, for example, in German there are more compound words than that in English. So when

translating cross languages, there are a plenty of cases that a single word will be translated to multiple words and multiple words correspond to a single conversely. We focus on such a case: from a German sentence: "ich höre zu" to "i'm listening". A very frequent word "zu" which corresponds to "to" in English, is dropped from the sentence. The design of reordering noise is as followed:

1. For each position  $i$ , sample a probability  $p_i \sim \text{Uniform}(0, 1)$
2. If  $p_i < p_{ins}$ , sample a word  $e$  from the most frequent  $V_{ins}$  target words and insert it before the position  $i$

We limit the insertion word in a set consisting of the top frequent word in the target language  $V_{ins}$

### Deletion Noise

The deletion noise is just a contrary case of insertion noise. Because we are limited to generate only ne word per source word, it is also possible that a target word in the reference is not related to any source word. For example for "eine der besten" the corresponding translation is "one of the best". We need to add an extra preposition in the target sentence. To simulate such situation, we drop some words randomly from a clean target sentence.

1. For each position  $i$ , sample a probability  $p_i \sim \text{Uniform}(0, 1)$
2. If  $p_i < p_{del}$ , drop the word in the position  $i$

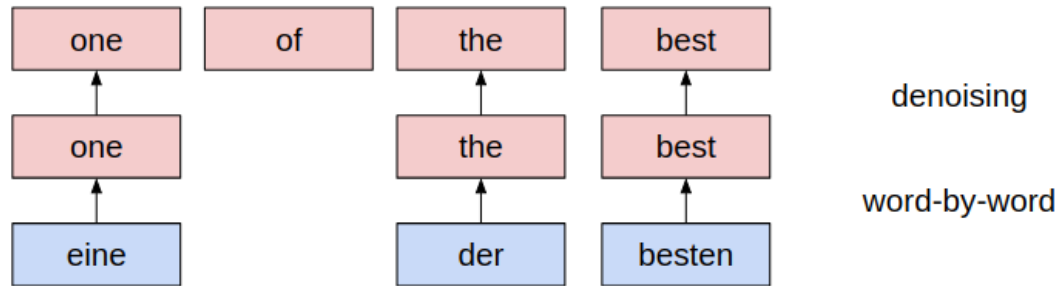


Figure 5.3: Deletion noise





# Chapter 6

## Experiments

### 6.1 Corpus Statistics

### 6.2 title

### 6.3 title

	Train		German	English	French
		Sentences	100M	100M	100M
		Running Words	1880M	2360M	3017M
		Vocabulary	1254k	523k	660k
Test		newstest2016		newstest2014	
		German	English	French	English
	Sentences	2999	2999	3003	3003
	Running Words	62506	64619	81165	71290
	Vocabulary Size	11978	8645	10899	9200
	OOV Rates	4116 (6.6%)	1643 (2.5%)	1731 (2.1%)	1299 (1.8%)
	LM perplexity	211.0	109.6	51.2	84.6

Search vocabulary in testing: 50k (src/tgt)

Table 6.1: Translation results on German↔English newstest2016 and French↔English newstest2014.

System	de-en Bleu [%]	en-de Bleu [%]	fr-en Bleu [%]	en-fr Bleu [%]
Word-by-Word	11.1	6.7	10.6	7.8
+ LM (5-gram) + tgt w/ high LM score for OOV	12.9	8.9	12.7	10.0
+ LM (5-gram) + copy from src for OOV	14.5	9.9	13.6	10.9
+ Denoising (RNN)	16.2	10.6	15.8	13.3
+ Denoising (Transformer)	17.2	11.0	16.5	13.9
Lample et al. [2017]	13.3	9.6	14.3	15.1
Artetxe et al. [2017b]	-	-	15.6	15.1

Table 6.2: Word-by-word translation from German to English

	Accuracy [%]	Bleu [%]
5M	44.9	9.7
10M	51.6	10.1
50M	59.4	10.8
100M	61.2	11.2

	Vocabulary	Bleu [%]
	Merges	
	20k	10.4
BPE	50k	12.5
	100k	13.0
	Cross-lingual training	
	20k	14.4
Word	50k	14.4
	100k	14.5
	200k	14.4

$d_{\text{per}}$	$p_{\text{del}}$	$p_{\text{ins}}$	$V_{\text{ins}}$	Bleu [%]
2				14.7
3				14.9
5				14.9
3	0.1			15.7
	0.3			15.1
			10	16.8
			50	17.2
3	0.1	0.1	500	16.8
			5000	16.5

Vocabulary			No LM Bleu [%]	With LM Bleu [%]	Denoising Bleu [%]
Word			11.2	14.5	17.2
Mikolov et al. [2013c]	threshold	100	11.1	13.7	15.6
		500	11.0	13.7	16.2
		2000	10.7	14.0	16.5
Top frequent	count	50k	12.0	15.7	16.8

Table 6.2: Word embedding vocabulary cut-off

Table 6.3: Phrase embedding vocabulary cut-off

Bleu [%]	20k	50k	100k
50k	11.1	11.3	11.2
100k	11.2	11.2	11.1
150k	10.9	10.9	-

Bleu [%]	50k	100k	150k
50k	11.3	-	-
100k	11.9	11.9	-
150k	12.0	11.9	11.9
200k	12.0	-	-



## Chapter 7

### Conclusion

For unsupervised machine translation system, the context-aware beam search language model help at the lexicon choice step.

The denoising networks which aimed at the insertion/deletion/reordering noise in the word-by-word translation sentence works well for fertility and localized reordering problem. Even though BPE is known to be an effective way to overcome the rare word problem in standard NMT, BPE embedding performs worse than word embedding in our case especially when vocabulary is small.

Word-by-word translation based on cross-lingual word embedding depends highly on the frequent word mappings. We found that phrase embedding only helps in word-by-word translation with context-aware beam search, it works not as good as that under the processing of denoising autoencoder.



## **Appendix A**

### **Appendix**





## List of Figures

2.1	Neural machine translation – example of a deep recurrent architecture (Luong et al. [2015]) . . . . .	11
2.2	Global attention model (Luong et al. [2015]) . . . . .	13
2.3	Local attention model (Luong et al. [2015]) . . . . .	14
2.4	The Transformer model architecture (Vaswani et al. [2017]) . . . . .	16
2.5	Illustration of unsupervised machine translation: A) the two original monolingual data; B) Initialization, the two distribution are roughly aligned; C) Denoising autoencoder, make the distribution closer to normal one in corresponding language D) Back-translation, improve the both translation model in iterative way (Lample et al. [2018]) . .	19
3.1	Global attention model (Mikolov et al. [2013a]) . . . . .	23
3.2	A cross-lingual embedding space between German and English (Ruder et al. [2017]) . . . . .	26
5.1	Reordering noise . . . . .	37
5.2	Insertion noise . . . . .	38
5.3	Deletion noise . . . . .	39



## List of Tables

6.1	Translation results on German $\leftrightarrow$ English <b>newstest2016</b> and French $\leftrightarrow$ English <b>newstest2014</b> . . . . .	42
6.2	Word-by-word translation from German to English . . . . .	42
6.2	Word embedding vocabulary cut-off . . . . .	43
6.3	Phrase embedding vocabulary cut-off . . . . .	43



## Bibliography

- Oliver Adams, Adam Makarucha, Graham Neubig, Steven Bird, and Trevor Cohn. Cross-lingual word embeddings for low-resource language modeling. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 1, Long Papers*, volume 1, pages 937–947, 2017.
- Mikel Artetxe, Gorka Labaka, and Eneko Agirre. Learning bilingual word embeddings with (almost) no bilingual data. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, volume 1, pages 451–462, 2017a.
- Mikel Artetxe, Gorka Labaka, Eneko Agirre, and Kyunghyun Cho. Unsupervised neural machine translation. *arXiv preprint arXiv:1710.11041*, 2017b.
- Piotr Bojanowski, Edouard Grave, Armand Joulin, and Tomas Mikolov. Enriching word vectors with subword information. *arXiv preprint arXiv:1607.04606*, 2016.
- Hailong Cao, Tiejun Zhao, Shu Zhang, and Yao Meng. A distribution-based model to learn bilingual word embeddings. In *Proceedings of COLING 2016, the 26th International Conference on Computational Linguistics: Technical Papers*, pages 1818–1827, 2016.
- Yong Cheng, Wei Xu, Zhongjun He, Wei He, Hua Wu, Maosong Sun, and Yang Liu. Semi-supervised learning for neural machine translation. *arXiv preprint arXiv:1606.04596*, 2016.
- Trevor Cohn and Mirella Lapata. Machine translation by triangulation: Making effective use of multi-parallel corpora. In *Proceedings of the 45th Annual Meeting of the Association of Computational Linguistics*, pages 728–735, 2007.
- Alexis Conneau, Guillaume Lample, Marc’Aurelio Ranzato, Ludovic Denoyer, and Hervé Jégou. Word translation without parallel data. *arXiv preprint arXiv:1710.04087*, 2017.
- Long Duong, Hiroshi Kanayama, Tengfei Ma, Steven Bird, and Trevor Cohn. Learning crosslingual word embeddings without bilingual corpora. *arXiv preprint arXiv:1606.09403*, 2016.

- Jonas Gehring, Michael Auli, David Grangier, Denis Yarats, and Yann N Dauphin. Convolutional sequence to sequence learning. *arXiv preprint arXiv:1705.03122*, 2017.
- Di He, Yingce Xia, Tao Qin, Liwei Wang, Nenghai Yu, Tieyan Liu, and Wei-Ying Ma. Dual learning for machine translation. In *Advances in Neural Information Processing Systems*, pages 820–828, 2016.
- Yedid Hoshen and Lior Wolf. An iterative closest point method for unsupervised word translation. *CoRR*, abs/1801.06126, 2018. URL <http://arxiv.org/abs/1801.06126>.
- Yunsu Kim, Julian Schamper, and Hermann Ney. Unsupervised training for large vocabulary translation using sparse lexicon and word classes. In *Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: Volume 2, Short Papers*, volume 2, pages 650–656, 2017.
- Guillaume Lample, Ludovic Denoyer, and Marc’Aurelio Ranzato. Unsupervised machine translation using monolingual corpora only. *arXiv preprint arXiv:1711.00043*, 2017.
- Guillaume Lample, Myle Ott, Alexis Conneau, Ludovic Denoyer, and Marc’Aurelio Ranzato. Phrase-based & neural unsupervised machine translation. *arXiv preprint arXiv:1804.07755*, 2018.
- Minh-Thang Luong, Hieu Pham, and Christopher D Manning. Effective approaches to attention-based neural machine translation. *arXiv preprint arXiv:1508.04025*, 2015.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*, 2013a.
- Tomas Mikolov, Quoc V Le, and Ilya Sutskever. Exploiting similarities among languages for machine translation. *arXiv preprint arXiv:1309.4168*, 2013b.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg S Corrado, and Jeff Dean. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, pages 3111–3119, 2013c.
- Malte Nuhn and Hermann Ney. Em decipherment for large vocabularies. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, volume 2, pages 759–764, 2014.

- Malte Nuhn, Arne Mauser, and Hermann Ney. Deciphering foreign language by combining language models and context vectors. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics: Long Papers-Volume 1*, pages 156–164. Association for Computational Linguistics, 2012.
- Jeffrey Pennington, Richard Socher, and Christopher Manning. Glove: Global vectors for word representation. In *Proceedings of the 2014 conference on empirical methods in natural language processing (EMNLP)*, pages 1532–1543, 2014.
- Sujith Ravi and Kevin Knight. Deciphering foreign language. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies-Volume 1*, pages 12–21. Association for Computational Linguistics, 2011.
- Sebastian Ruder, Ivan Vulić, and Anders Søgaard. A survey of cross-lingual word embedding models. *arXiv preprint arXiv:1706.04902*, 2017.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, pages 5998–6008, 2017.
- Ivan Vulic and Anna-Leena Korhonen. On the role of seed lexicons in learning bilingual word embeddings. 2016.
- Chao Xing, Dong Wang, Chao Liu, and Yiye Lin. Normalized word embedding and orthogonal transform for bilingual word translation. In *Proceedings of the 2015 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1006–1011, 2015.
- Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron Courville, Ruslan Salakhudinov, Rich Zemel, and Yoshua Bengio. Show, attend and tell: Neural image caption generation with visual attention. In *International conference on machine learning*, pages 2048–2057, 2015.
- Meng Zhang, Yang Liu, Huanbo Luan, and Maosong Sun. Adversarial training for unsupervised bilingual lexicon induction. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, volume 1, pages 1959–1970, 2017.